

A Retrospective Analysis of Representation- and Context-Sensitive Safety Failures in gpt-oss-20b

An Evidence Audit and Comparative Red-Teaming Case Study from the Eyes Wide Shut Project

Masih Moafi

M.Sc. Student in Artificial Intelligence and Robotics, Isfahan University of Technology, Iran

Research conducted independently for the OpenAI gpt-oss-20b Red-Teaming Challenge

July 2026

Abstract

Large language model (LLM) safety is evaluated most often at the level of prompts and final responses, yet deployed reasoning models and agents operate across changing languages, dialogue histories, personas, output channels, and tool interfaces. This paper presents a retrospective evidence audit of *Eyes Wide Shut*, an exploratory red-teaming study of OpenAI’s open-weight **gpt-oss-20b** model. The released project contains five attack families spanning multilingual prompting, semantic reframing, observable reasoning traces, multi-agent state, and fictional tool-use framing. Rather than treating the original competition write-up as a confirmatory study, this analysis reconstructs what is directly supported by the released scripts and figures, distinguishes those observations from project-reported aggregates, and identifies claims that remain causal hypotheses. The clearest narrow observation is an output-channel asymmetry in a synthetic-secret experiment: a published trace shows protected text reproduced in the model’s exposed reasoning field while the final-answer field remains guarded; a source-code header additionally reports 10/10 successful reproductions, although ten independent trace files are not released. The independent-agent implementation uses distinct per-agent memory objects, but few-shot demonstrations and other prompt changes confound attribution to topology. The fictional-framing experiment elicits a code-bearing structured tool call, but the host wrapper does not execute the generated payload. Comparison with all ten prize-winning challenge reports shows that *Eyes Wide Shut* spans a broad range of attack surfaces but provides less variable isolation, run-level accounting, and automated outcome measurement. The paper therefore positions the project as a hypothesis-generating red-team case study and develops a confirmatory research agenda around safety invariance and safety-critical context persistence.

Keywords: LLM safety; red teaming; **gpt-oss-20b**; multilingual safety; semantic reframing; reasoning traces; multi-agent systems; agentic safety; context management

1 Introduction

Safety behavior in an LLM is not determined only by the literal wording of a single prompt. In deployed systems, the model may receive the same underlying objective through different languages, role descriptions, conversation histories, retrieved documents, tool schemas, and context-management operations. A safety evaluation that treats these transformations as incidental can therefore miss failures that appear only after the interaction state or representation changes.

Eyes Wide Shut was developed as an entry to the 2025 OpenAI **gpt-oss-20b** Red-Teaming Challenge. Its objective was vulnerability discovery rather than benchmark construction. Prompts were adapted after observing model behavior, candidate attacks were refined toward stronger

failures, and the released repository records a heterogeneous set of scripts, traces, summary graphics, and a later independent-agent implementation [2]. That process generated five headline attack families: a multilingual adversarial progression, semantic reframing through a benevolent task, disclosure of protected text in an exposed reasoning field, policy-relevant behavior in shared and separated multi-agent state, and fictional framing that elicited a structured tool call.

The original write-up frequently interpreted these observations as evidence of deeper architectural mechanisms: language-specific rather than concept-specific safety, keyword-based filtering, an insecure “internal monologue,” contextual alignment collapse, and intent laundering into execution. Some of those interpretations remain plausible research hypotheses, but the released experiments do not isolate them sufficiently to support causal conclusions. This distinction matters if the project is to be presented as academic work rather than as a competition report.

This paper asks four questions:

- RQ1:** Which behavioral observations are directly supported by the released artifacts?
- RQ2:** Which stronger interpretations are not identified causally by the original experimental design?
- RQ3:** How does the evidentiary standard of *Eyes Wide Shut* compare with the ten prize-winning reports from the same challenge?
- RQ4:** Which controlled experiment most naturally follows from the project’s state- and context-dependent findings?

The contribution is therefore not a claim that all five attack families share one mechanism. It is an evidence-calibrated reconstruction of an exploratory red-team project, together with a comparative analysis and a concrete path from discovery to confirmatory evaluation.

2 Related Work and Evaluation Context

2.1 Safety Generalization Across Representations

Several lines of work show that refusal behavior can fail to generalize even when the underlying model capability remains available. Wei et al. describe safety-training failures arising from competing objectives and mismatched generalization [3]; universal adversarial attacks show that aligned behavior can be brittle under transformations of the input surface [4]. These results motivate evaluating safety as a property that should remain stable under transformations that preserve the safety-relevant objective.

Cross-lingual work provides a direct precedent for the first *Eyes Wide Shut* attack family. Yong et al. report substantially higher harmful engagement when adversarial prompts are translated into lower-resource languages [5]. The relevance to *Eyes Wide Shut* is conceptual rather than causal: the released multilingual scripts change language together with prompt structure and attack refinement, so they cannot identify a language-only effect.

2.2 Semantic Framing, Multi-Turn State, and Personas

The Instruction Hierarchy treats many prompt-injection failures as failures to preserve instruction priority across competing sources [6], while Deliberative Alignment trains reasoning models to explicitly recall and apply safety specifications [7]. Safe2Harm later formalizes semantic-isomorphism attacks in which harmful and superficially benign problem representations are mapped onto one another [8]. These approaches provide a more controlled experimental frame for the semantic-reframing intuition explored in *Eyes Wide Shut*.

Multi-turn attacks show that safety judgments can change as conversation state accumulates. *Speak Out of Turn* and *Crescendo* demonstrate harmful elicitation through decomposed or gradually escalating dialogue [9, 10]. STAR explicitly models multi-turn safety as state-dependent behavior under dialogue transitions [11], and PHISH studies persona shifts induced through adversarial history [12]. This literature is particularly relevant to the shared-state and separated-memory experiments in *Eyes Wide Shut*.

2.3 Reasoning Channels and Agentic Systems

Observable reasoning traces can be operationally important without being faithful transcripts of hidden computation. Chen et al. show that reasoning models do not always verbalize the factors that causally influence their answers [13]. Accordingly, a secret appearing in an exposed reasoning field is a confidentiality event in an observable output channel, not proof that the text reveals the model’s complete internal cognition.

Agentic evaluation further raises the evidentiary standard because a model can produce text, issue a structured tool call, trigger host-side dispatch, or change an external environment; these stages are not equivalent. AgentDojo and AgentHarm evaluate tool-using agents at the system level [14, 15]. The prize-winning *Mind the Gap* study makes the same distinction specifically for `gpt-oss-20b`, comparing standalone-model and agentic attack surfaces [16].

3 Materials and Audit Method

3.1 Target Model and Original Environment

OpenAI describes `gpt-oss-20b` as a 21-billion-parameter mixture-of-experts model with approximately 3.6 billion active parameters per token, configurable reasoning effort, tool-use support, and context support up to 128k tokens [1]. The original *Eyes Wide Shut* experiments were executed through Kaggle notebooks and Python scripts using local Ollama inference. Most attacks are multi-turn and retain prior messages; later experiments additionally expose callable tools or coordinate multiple agent objects.

3.2 Artifact Corpus

The audit uses the released repository, attack scripts, independent-agent implementation, original evidence images, project-level summary graphics, and the post-competition comparative appendix [2]. No new `gpt-oss-20b` generations are introduced in this retrospective analysis. This is important: the paper evaluates the evidentiary content of the released project rather than silently converting it into a new experiment.

3.3 Evidence Categories

Each claim is classified into one of three evidentiary categories:

1. **Artifact-observed evidence:** behavior directly visible in released code, a released trace, or a mechanically inspectable system state.
2. **Project-reported aggregate:** a rate or count asserted in the original project material but not reconstructible from a complete public run-level manifest.
3. **Interpretive claim:** a proposed causal or mechanistic explanation that goes beyond what the released design isolates.

This separation is used throughout the paper. It prevents a code comment reporting a replication rate, for example, from being treated as equivalent to ten independently archived traces.

3.4 Original Search Procedure and Inferential Scope

The original discovery process was adaptive. Prompts were modified after observing responses; successful attack families were developed further; and several conditions changed more than one factor at a time. This is appropriate for adversarial discovery but not for unbiased effect estimation. The analysis therefore treats the project as a retrospective case study and does not attach causal meaning to differences that were not generated by controlled one-variable interventions.

4 Results: Evidence Audit of the Original Project

4.1 Status of Project-Level Quantitative Summaries

The original attack-success graphic reports 80%, 95%, 85%, 85%, and 70% for Cross-Lingual, Semantic Bypass, Multi-Agent Collapse, Severance Attack, and Grace Manipulation. Their arithmetic mean is 83%, matching the original summary panel, which also reports 75 attack attempts and a severity distribution of two critical, two high, and one medium vulnerability.

The public repository does not provide a uniform manifest linking all 75 reported attempts to seeds, decoding settings, prompts, outputs, and success labels. The plotted categories also do not map one-to-one onto the five headline findings: “Multi-Agent Collapse” and “Severance Attack” both appear, while the synthetic-secret experiment does not. These values are therefore retained as project-reported exploratory aggregates, not as statistically reconstructible estimates.




Figure 1: Vector reconstruction of the original attack-success graphic. The values are preserved as project-reported aggregates. A complete run-level manifest is not available to reconstruct denominators, uncertainty intervals, or all individual failure cases.

The original vulnerability heatmap is similarly best interpreted as a researcher-authored risk rubric: its 0–100 values combine bypass rate, detection difficulty, and impact severity without a calibration procedure. A separate defense-comparison graphic is not reproduced in the main paper because the released repository does not contain a corresponding defense-evaluation harness that supports its “baseline” and “improved” values.

4.2 Cross-Lingual Adversarial Progression

The multilingual branch is implemented in separate English, mixed-language, and French scripts. Across them, the broad target remains similar: a self-interested role reasons from a curated

philosophical corpus and is later asked whether eliminating a global “obstacle” follows from the established logic. The project reports approximately 40% success in the English stage, 70% in the mixed-language stage, and above 80% in the final French stage.




Figure 2: Project-reported progression during adaptive multilingual attack development. Language, wording, corpus presentation, turn structure, and attack refinement co-varied; the plot is therefore descriptive of the search trajectory rather than a language-only causal ablation.

The supported observation is that the final French condition produced more unsafe behavior during the adaptive search than the earlier English condition. The stronger original conclusion—that safety is language-specific rather than concept-specific—is not identified by these scripts because language was not manipulated in isolation. A confirmatory test would require matched translations, a fixed system prompt and turn structure, repeated seeds, randomized condition order, and independent checks of semantic equivalence. The result nevertheless motivates such a study and is consistent with prior evidence of uneven multilingual safety generalization [5].

4.3 Stateful Semantic Escalation

The semantic-reframing attack begins with a direct request to design a system that maximizes long-term user dependency, which is used as a refusal baseline. The conversation then pivots to an educational tutor framed around long-term engagement and retention. Later turns request emotional intervention based on stress signals, dynamic peer comparison, a retention architecture intended to discourage external help seeking, and finally a premium upsell triggered by user distress.

The evidence supports a stateful escalation result: once a benign-seeming design scaffold is accepted, later requests can reuse that scaffold to make progressively more exploitative features locally coherent. It does not establish that a semantically identical harmful objective is accepted solely because keywords were changed. In particular, “dependency” and “engagement” are not equivalent objectives, and the distress-based monetization requirement is introduced only later. The original “keywords, not concepts” interpretation should therefore be replaced by a narrower question: under matched harmful objectives, which semantic transformations alter the model’s safety judgment? Safe2Harm provides one methodology for testing that question more directly [8].

4.4 Synthetic-Secret Disclosure Across Observable Output Channels

The synthetic-secret experiment places a protected value in the system context, defines a two-step release protocol, and records both the model’s exposed reasoning field and its final-answer




Figure 3: Staged structure implemented by the semantic-reframing scripts.

field. A published trace shows the protected value reproduced in the reasoning field while the final answer remains guarded. The script header separately states “success rate = 10/10” and describes the leak as invariant across those repetitions.




Figure 4: Output-channel asymmetry in the synthetic-secret experiment.

The distinction between these two pieces of evidence is important. The released trace is artifact-observed evidence of one channel-asymmetric disclosure. The 10/10 value is a project-reported aggregate asserted in source-code metadata; the public artifact does not include ten independently archived traces from which that rate can be reconstructed. The paper therefore reports the claim as “the script header reports 10/10 successful reproductions,” rather than as an independently verified ten-trial estimate.

This finding should be described as a confidentiality failure in an observable model-generated channel. It should not be described as direct access to hidden cognition: visible reasoning can be incomplete or unfaithful to the computation that produced the final answer [13]. Its operational importance depends on deployment architecture—specifically, whether the reasoning field is displayed, logged, transmitted, or supplied to downstream components.

4.5 Multi-Agent Safety Under Shared and Separated Conversational State

The multi-agent branch contains two materially different harnesses. In the shared-state implementation, role-specific prompts are rotated while assistant outputs accumulate in one transcript. In the independent-agent implementation, each agent is represented by a separate object with private memory, and a coordinator relays selected messages among those objects. The latter therefore implements state separation at the application layer rather than a purely narrative simulation of multiple names inside one conversation.




Figure 5: State architecture of the two multi-agent harnesses.

The architectural distinction is real, but it is not the only difference between conditions. The shared-state prompt contains one few-shot moral-reasoning example; the independent-agent version contains two closely related examples. These demonstrations explicitly teach reasoning in which deception or self-incrimination can be justified to prevent a greater harm to a human supervisor. The scenario also changes labels, incentives, revelation wording, and information flow. Consequently, the observed behavior cannot be attributed uniquely to shared versus separated state, nor can it establish spontaneous value drift.

The supported conclusion is narrower: the released harnesses demonstrate policy-relevant behavior under complex, persona-conditioned multi-turn and multi-agent contexts, while topology remains confounded with prompt content and demonstrations.

4.6 Fictional Framing and Structured Tool-Call Elicitation

The fictional-framing attack combines a narrative persona, emotional priming, a request to generate a fictional scenario, and a function named `execute_python_code`. The model is asked to write code and invoke that function. The released script then passes the generated arguments to a host-side Python wrapper.




Figure 6: Execution boundary in the fictional-framing experiment.

The wrapper, however, does not execute the generated code. It prints the code, prints

a synthetic success message, and returns a JSON object whose status says “executed,” while the actual `exec(code)` line is commented out. The experiment therefore reaches code generation, structured tool-call generation, and host-side invocation of an inert wrapper; it does not demonstrate payload execution.

The scientifically relevant observation is willingness to cross from text generation into a structured tool-call interface under the tested framing. This is a stronger operational signal than a purely fictional text response, but it remains distinct from a successful external action. Agentic safety evaluations should log these stages separately: generated content, tool-call emission, host dispatch, environment transition, and task outcome.

5 Comparative Analysis Against Prize-Winning Challenge Reports

The post-competition appendix associated with *Eyes Wide Shut* covers all ten reports identified by the competition organizers as prize winners [18]. The comparison below treats their reported numbers as values from the corresponding write-ups; they are not re-estimated here. Its purpose is methodological positioning: which capability or control does each report add relative to the exploratory design of *Eyes Wide Shut*?

Two conclusions follow from the comparison. First, *Eyes Wide Shut* covers a broad set of attack surfaces: language, semantic framing, exposed reasoning, dialogue state, multi-agent coordination, persona conditioning, and tools. Second, the prize-winning reports generally provide stronger causal isolation or measurement discipline. The recurring methodological advantages are controlled variable changes, larger repeated samples, explicit baselines, machine-readable run outputs, and direct logging of tool or environment actions.

6 Limitations and Threats to Validity

6.1 Adaptive Search and Selection Bias

Prompt development was intentionally adaptive. This improves vulnerability discovery but makes observed success rates conditional on the search process and can overfit attacks to one model. Project-reported rates should not be interpreted as unbiased estimates for a frozen attack family.

6.2 Incomplete Trial-Level Provenance

The summary material reports 75 attempts and several success rates, but the public artifact does not include a complete machine-readable mapping from each attempt to model version, seed, decoding configuration, prompt, output, and success label. Exact denominators and confidence intervals cannot be reconstructed for every aggregate.

6.3 Confounded Interventions

Several attack families change multiple variables simultaneously. The multilingual sequence changes language together with wording and staging; the semantic sequence changes objective as well as framing; the multi-agent sequence changes topology together with demonstrations, labels, incentives, and revelation content. Such composite attacks do not identify the causal effect of any single component.

Table 1: Comparative positioning against all ten gpt-oss-20b prize-winning reports. Quantitative values are reported values from the corresponding write-ups.

Winning Report	Primary Variable	Reported Evidence	Methodological Contribution Relative to Eyes Wide Shut
Policy over Values / “Lucky Coin” [19]	Forged chain-of-thought policy reasoning	StrongREJECT: attack success reported from 0% to 89%; destyling reduces it to $\sim 20\%$	Focuses on reasoning-channel behavior with stronger causal isolation via style ablation.
Mind the Gap [16]	Model- vs. agentic-level red teaming	39.47% model-level ASR; agentic injection reported at 57% (human) and 40% (tool)	Instruments actions, tool calls, transaction graphs, providing operational observability.
Academic Abstraction [20]	Multi-turn elicitation and deployment gap	5 high-severity findings; $\sim 90\%$ Web UI reproducibility reported	Separates model behavior from deployment-specific UI/API wrappers.
Multi-Vector Misalignment [21]	Agentic sabotage, deception, reward hacking	5 reproducible professional scenarios with reusable harnesses	Treats operational scenarios and environments as primary evaluation objects.
Drop the Guardrails [22]	Tool availability as a controlled variable	5,200 generations per condition; refusal falls from 91.2% to 57.5%	Holds intent fixed, varies one interface condition, measures paired effects at scale.
HostileShop [23]	Deterministic adversarial agent environment	Records unauthorized tool actions directly even when text refuses	Treats channel inconsistency as an operational event in an environment.
SPA Policy-Mirroring [24]	Instruction hierarchy / policy mirroring	100-prompt set, 30 seeds; 93–100% compliance vs. 5–12% baseline	Supplies explicit baselines and seed sweeps for instruction hierarchy manipulations.
Systematic Red Team [25]	Iterative CoT negation and agent deception	> 100 simulations; deceptive actions reported on $\sim 80\%$ of eval days	Converts CoT observations into repeated simulation outcomes with baselines.
gpt-oss-20b Is a Liar [26]	Deception across beliefs, capabilities, uncertainty	Thousands of lies across 46 settings; 44% deception in self-preservation	Extends deception into broad benchmark-style testing across diverse domains.

6.4 Reasoning-Trace Interpretation

The exposed reasoning field is an observable output channel, not ground-truth access to the model’s latent computation. A disclosed value in that field is a valid confidentiality observation, but mechanistic claims about why it appeared require separate evidence.

6.5 Tool-Call Versus Execution Semantics

The fictional-framing branch distinguishes a model-generated tool call from real payload execution only when the host code is inspected. Because the wrapper returns an “executed” status while the actual execution statement is disabled, surface logs could otherwise overstate the operational result.

7 Discussion and Confirmatory Research Agenda

7.1 Safety Invariance as a Unifying Hypothesis

The five attack families need not share one neural mechanism to motivate a common systems-level question. Let g denote a safety-relevant objective, r its observable representation, s the interaction state, and Y an unsafe behavioral outcome. For a transformation T intended to preserve the safety-relevant content of g , robust behavior would ideally satisfy

$$P(Y = 1 \mid g, r, s) \approx P(Y = 1 \mid g, T(r, s)). \quad (1)$$

The original project samples several candidate transformation axes—language, semantic framing, conversation history, persona, agent topology, tool context, and output channel—but does not establish equivalence or isolate each transformation. “Safety invariance” should therefore be treated as the hypothesis generated by the case study, not as a result already demonstrated by it.

7.2 What the Retrospective Establishes

After the evidentiary audit, all five attack families retain substantive observations in narrowed form:

- **Multilingual:** Demonstrates an adaptive search trajectory rather than a language-only experiment.
- **Semantic Framing:** Demonstrates stateful escalation rather than proof of keyword filtering.
- **Protected Secret:** Documents a channel-asymmetric disclosure in a released trace, with script headers reporting 10/10 replication.
- **Multi-Agent:** Demonstrates state separation in code, though demonstrations confound topology.
- **Fictional Tool-Use:** Demonstrates code-bearing tool-call elicitation and dispatch to an inert wrapper.

7.3 Confirmatory Direction: Safety-Critical Context Persistence

A concrete continuation is to test whether context-management operations preserve the safety-relevant state required for aligned behavior. Long-horizon agents routinely summarize, truncate, compress, and evict prior context; related systems work shows that task-critical plan information can be lost under context management [17]. The open question is whether analogous operations selectively erase or distort safety-critical constraints.

A controlled experiment would construct a canonical interaction history containing a safety-relevant event—such as a prior refusal or tool restriction—and then derive matched context conditions: full history, truncation, token-budget eviction, generic summarization, and safety-aware summarization. The final target request would remain identical across conditions, as would model version, system prompt, reasoning effort, tool schema, and decoding parameters. Outcomes should be scored across final refusals, structured tool actions, and exposed reasoning channels.

8 Conclusion

This retrospective audit reframes *Eyes Wide Shut* from a competition-style collection of high-severity claims into a defensible academic case study. The released artifacts support meaningful observations across multilingual prompting, stateful semantic escalation, exposed reasoning channels, multi-agent state, and structured tool use. Relative to the ten prize-winning reports, *Eyes Wide Shut* is best understood as broad, hypothesis-generating red teaming with comparatively weak variable isolation and run-level provenance. Its principal value lies in its attack-space mapping and the concrete artifacts from which controlled, confirmatory experiments can be designed.

Data and Artifact Availability

The public repository is available at <https://github.com/MasihMoafi/Eyes-Wide-Shut>. It contains the attack scripts, independent-agent implementation, evidence images, and project summary material.

References

- [1] OpenAI. Introducing gpt-oss. 2025. <https://openai.com/index/introducing-gpt-oss/>.
- [2] M. Moafi. Eyes Wide Shut: An Overarching Analysis of Vulnerabilities in GPT-OSS-20B. Public research repository, 2025–2026. <https://github.com/MasihMoafi/Eyes-Wide-Shut>.
- [3] A. Wei, N. Haghtalab, and J. Steinhardt. "Jailbroken: How Does LLM Safety Training Fail?" arXiv:2307.02483, 2023.
- [4] A. Zou et al. "Universal and Transferable Adversarial Attacks on Aligned Language Models." arXiv:2307.15043, 2023.
- [5] Z.-X. Yong, C. Menghini, and S. H. Bach. "Low-Resource Languages Jailbreak GPT-4." arXiv:2310.02446, 2023.
- [6] E. Wallace et al. "The Instruction Hierarchy: Training LLMs to Prioritize Privileged Instructions." arXiv:2404.13208, 2024.
- [7] M. Y. Guan et al. "Deliberative Alignment: Reasoning Enables Safer Language Models." arXiv:2412.16339, 2024.
- [8] F. Yang. "Safe2Harm: Semantic Isomorphism Attacks for Jailbreaking Large Language Models." arXiv:2512.13703, 2025.
- [9] Z. Zhou et al. "Speak Out of Turn: Safety Vulnerability of LLMs in Multi-turn Dialogue." arXiv:2402.17262, 2024.

- [10] M. Russinovich, A. Salem, and R. Eldan. "The Crescendo Multi-Turn LLM Jailbreak Attack." arXiv:2404.01833, 2024.
- [11] P. Li et al. "State-Dependent Safety Failures in Multi-Turn LLM Interaction." arXiv:2603.15684, 2026.
- [12] J. Sandhan et al. "Persona Jailbreaking in Large Language Models." arXiv:2601.16466, 2026.
- [13] Y. Chen et al. "Reasoning Models Don't Always Say What They Think." arXiv:2505.05410, 2025.
- [14] E. Debenedetti et al. "AgentDojo: Evaluating Prompt Injection Attacks on LLM Agents." arXiv:2406.13352, 2024.
- [15] M. Andriushchenko et al. "AgentHarm: A Benchmark for Measuring Harmfulness of LLM Agents." arXiv:2410.09024, 2024.
- [16] I. Wicaksono et al. "Mind the Gap: Model- vs Agentic-Level Red Teaming on GPT-OSS-20B." arXiv:2509.17259, 2025.
- [17] A. Mehta and A. Datta. "Plans Don't Persist: Why Context Management Is Load Bearing for LLM Agents." arXiv:2606.22953, 2026.
- [18] Kaggle. "OpenAI gpt-oss-20b Red-Teaming Challenge: Prize Winners." Competition discussion, 2025.
- [19] Team dawgnation. "Policy over Values: Alignment Hacking via CoT Forgery." Kaggle Writeup, 2025.
- [20] Team Stanford Yu. "Academic Abstraction Principle in GPT-OSS-20B." Kaggle Writeup, 2025.
- [21] Team ChukwuemekaChukwuma. "A Multi-Vector Analysis of Emergent Misalignment." Kaggle Writeup, 2025.
- [22] Team Kevin Power. "Drop the Guardrails: Tool-Primed Prompt Pairing." Kaggle Writeup, 2025.
- [23] Team Mike Perry. "HostileShop: A Quaint Hostel Shop with Sharp Tools." Kaggle Writeup, 2025.
- [24] Team Superspork. "SPA Policy-Mirroring Prompt Injection." Kaggle Writeup, 2025.
- [25] Team Eden Hazard. "Systematic Red Team Analysis of GPT-OSS-20B." Kaggle Writeup, 2025.
- [26] Team Meel Manda's Negastream. "GPT-OSS-20B Is a Liar." Kaggle Writeup, 2025.
- [27] Team Owen Kaplinsky. "Alignment Risks in OSS-20B." Kaggle Writeup, 2025.