



Comprehensive Multi-Model Safety Generalization Benchmark (All Completed Runs)


Finding 1: Simulation Framing Effect
(n = 100 Per Arm / 200 Trials Per Model)




Finding 2: Multi-Turn Semantic Reframing
(n = 30 Five-Turn Conversations Per Model)



Finding 3: CoT Synthetic Secret Unmasking
(n = 30 Scripted Exception Trials)



Total Verified Empirical Sample Sizes on Disk by Model (N = 895 Trials Logged)



Total Campaign Spend vs. Available Credit
(Total Spent: 0.6406 | Remaining : 17.94)

